

Planning and Learning in Dynamical Systems (046194)

Homework 1, due 18/4/2019

Question 1 (Counting)

Given an integer number X , we would like to count in how many ways it can be represented as a sum of N natural numbers (with relevance to order) marked by $\Psi_N(X)$. For example, $\Psi_2(3) = 2$ since: $3 = 1 + 2 = 2 + 1$.

- Find the following: $\Psi_1(2), \Psi_2(4), \Psi_3(2), \Psi_N(N)$ and $\Psi_1(X)$.
- Find a recursive equation for $\Psi_N(X)$.
- Write a code in for finding $\Psi_N(X)$ using dynamic programming.
 - What is the time and memory complexity of your algorithm?
 - Find $\Psi_{12}(800)$.
- Now assume each natural number i is associated with some cost c_i . For a given X, N we are interested in finding the lowest cost combination of natural numbers $\{x_i\}_{i=1}^N$ satisfying $\sum_{i=1}^N x_i = X$.
 - Formulate the problem as a finite horizon decision problem: Define the state space, the action space and the cumulative cost function.
 - Bound the complexity of finding the best combination.

Question 2 (Language model)

In the city of Bokoboko the locals use a language with only 3 letters ('B','K','O'). After careful inspection of this language, researchers have reached two conclusions:

- Every word starts with the letter 'B'.
- Every consecutive letter is distributed only according to the previous letter as follows:

$$P(l_{t+1} | l_t) = \begin{matrix} & \begin{matrix} B & K & O & - \end{matrix} \\ \begin{matrix} B \\ K \\ O \end{matrix} & \begin{bmatrix} 0.1 & 0.325 & 0.25 & 0.325 \\ 0.4 & 0 & 0.4 & 0.2 \\ 0.2 & 0.2 & 0.2 & 0.4 \end{bmatrix} \end{matrix}$$

Where '-' represents the end of a word. For example, the probability of the word 'bko' is given by $0.325 \cdot 0.4 \cdot 0.4 = 0.052$.

- Find the probability of the following words: 'Bob', 'Koko', 'B', 'Bokk', 'Boooo'.
- We wish to find the most probable word in the language of length K .
 - Formulate the problem as a finite horizon decision problem: Define the state space, the action space and the **multiplicative** cost function.
 - Bound the complexity of finding the best combination.
 - Find a reduction from the given problem to an analogous problem with **additive** cost function instead.

- ### Question 3 (Path planning)

- a. Formulate the problem as a finite horizon decision problem: Define the state space, the action space and the cumulative cost function.
- b. What is the horizon of the problem?
- c. How many possible trajectories are there? How does the number of trajectories behaves as a function of N when $M = 2$? How does it behave as a function of N when $M = N$?
- d. Aharon, Moses's long lost war-buddy woke up confused next to Moses and decided to join him in his quest (needless to say, both mice suffer the same rare head injury).
 1. Explain what will happen if both mice ignore each other's existence and act 'optimal' with respect to the original problem.
 2. Assume both mice decided to coordinate their efforts and split the loot. How many states and actions are there now?
 3. Now their entire rarely-head-injured division has joined the journey. Assume there's a total of K mice, how many states and actions are there now?

Question 4 (MinMax dynamic programming)

In this problem we consider an adversarial version of finite horizon dynamic programming, which is suitable for solving 2-player games. In this setting, at time $k \in \{0, 1, 2, \dots, N-1\}$ the system is at state $s_k \in S_k$, the agent chooses an action $a_k \in A_k(s_k)$ according to the agent policy $\pi_k^a(s_k)$, and subsequently the opponent chooses an action b_k from the set of allowable opponent actions $B_k(s_k, a_k)$, according to the opponent policy $\pi_k^b(s_k, a_k)$.

The system then transitions to a new state according to:

$$s_{k+1} = f_k(s_k, a_k, b_k), \quad k = 0, 1, 2, \dots, N-1.$$

The instantaneous reward is denoted by $r(s_k, a_k, b_k)$, and, for an N-stage path $h_N = (s_0, a_0, b_0, \dots, s_{N-1}, a_{N-1}, b_{N-1}, s_N)$ the cumulative reward is

$$R_N(h_N) = \sum_{k=0}^{N-1} r_k(s_k, a_k, b_k) + r_N(s_N).$$

Given s_0 , the agent's goal is to find a policy π_a^* that maximizes the *worst-case* cumulative reward:

$$\pi_a^* \in \arg \max_{\pi_a} \min_{\pi_b} R_N(h_N)$$

- Formulate a dynamic programming algorithm for this problem. Explain what the value function represents in this case.
- What is the computational complexity of your algorithm?
- Could this approach be used to solve the game of tic-tac-toe? Explain what are the states, actions and rewards for this game.
- Could this approach be used to solve the game of chess? Explain.